

Next Steps Implementation Guide

DevOps Requirements - Detailed Implementation Plan

Based on the "next steps of the assignment.pdf" requirements, this document outlines exactly what needs to be implemented.

Overview of Requirements

The DevOps phase requires:

1. **Dockerfile** (✓ Done) - Expose port 8080, include .dockerignore
 2. **Environment Variables** (✓ Done) - MONGODB_URI, credentials, .env.example
 3. **Terraform** - Resource group, compute, networking for /rest accessibility
 4. **GitHub Actions** - Test on PR, deploy on push, build/push Docker image, terraform apply
 5. **Data Validation** - restaurants.json automation (✓ Done - skip invalid rows)
-

1. Terraform Infrastructure (TO DO)

1.1 Directory Structure

Create the following structure:

```
terraform/
├── main.tf      # Main configuration
├── variables.tf # Input variables
├── outputs.tf   # Output values
├── providers.tf # Provider configuration
├── container.tf # Container Instance resource
├── networking.tf # Virtual network, NSG
└── terraform.tfvars.example # Example variable values
```

1.2 providers.tf

hcl

```
terraform {
  required_version = ">= 1.5.0"

  required_providers {
    azurerm = {
      source  = "hashicorp/azurerm"
      version = "~> 3.80"
    }
  }

  # Local backend (per company clarification)
  backend "local" {
    path = "terraform.tfstate"
  }
}

provider "azurerm" {
  features {}

  # These will come from environment variables or GitHub Secrets:
  # ARM_SUBSCRIPTION_ID
  # ARM_TENANT_ID
  # ARM_CLIENT_ID
  # ARM_CLIENT_SECRET
}
```

1.3 variables.tf

hcl

```
variable "resource_group_name" {
  description = "Name of the Azure resource group"
  type       = string
  default    = "rg-restaurant-recommendation"
}

variable "location" {
  description = "Azure region"
  type       = string
  default    = "West Europe" # Or your preferred region
}

variable "container_name" {
  description = "Name of the container instance"
  type       = string
  default    = "restaurant-api"
}

variable "docker_image" {
  description = "Docker image to deploy"
  type       = string
  # Will be set by CI/CD: ghcr.io/<owner>/restaurant-recommendation:latest
}

variable "mongodb_uri" {
  description = "MongoDB Atlas connection string"
  type       = string
  sensitive  = true
}

variable "cpu_cores" {
  description = "CPU cores for container"
  type       = number
  default    = 1
}

variable "memory_gb" {
  description = "Memory in GB for container"
  type       = number
  default    = 1.5
}

variable "environment" {
  description = "Environment name (dev, staging, prod)"
  type       = string
  default    = "dev"
}
```

1.4 main.tf

hcl

```
# Resource Group
resource "azurerm_resource_group" "main" {
    name      = var.resource_group_name
    location  = var.location
```

```
    tags = {
        environment = var.environment
        project     = "restaurant-recommendation"
    }
}
```

Container Instance

```
resource "azurerm_container_group" "api" {
    name          = var.container_name
    location      = azurerm_resource_group.main.location
    resource_group_name = azurerm_resource_group.main.name
    os_type       = "Linux"
    ip_address_type = "Public"
    dns_name_label = "${var.container_name}-${var.environment}"
```

```
    container {
        name  = "restaurant-api"
        image = var.docker_image
        cpu   = var.cpu_cores
        memory = var.memory_gb
```

```
    ports {
        port  = 8080
        protocol = "TCP"
    }
```

```
    environment_variables = {
        "LOG_LEVEL"  = "INFO"
        "LOG_FORMAT" = "json"
    }
```

```
    secure_environment_variables = {
        "MONGODB_URI" = var.mongodb_uri
    }
```

```
    liveness_probe {
        http_get {
            path  = "/health"
            port  = 8080
            scheme = "Http"
        }
        initial_delay_seconds = 30
        period_seconds = 10
    }
```

```

period_seconds      = 10
failure_threshold   = 3
}

readiness_probe {
  http_get {
    path  = "/health"
    port  = 8080
    scheme = "Http"
  }
  initial_delay_seconds = 10
  period_seconds        = 5
}
}

tags = {
  environment = var.environment
  project     = "restaurant-recommendation"
}
}

```

1.5 outputs.tf

```

hcl

output "api_url" {
  description = "Public URL for the REST API"
  value       = "http://${azurerm_container_group.api fqdn}:8080"
}

output "api_fqdn" {
  description = "Fully qualified domain name"
  value       = azurerm_container_group.api.fqdn
}

output "api_ip_address" {
  description = "Public IP address"
  value       = azurerm_container_group.api.ip_address
}

output "resource_group_name" {
  description = "Resource group name"
  value       = azurerm_resource_group.main.name
}

output "rest_endpoint" {
  description = "The /rest endpoint URL"
  value       = "http://${azurerm_container_group.api fqdn}:8080/rest"
}

```

1.6 terraform.tfvars.example

hcl

```
# Copy to terraform.tfvars and fill in values
# DO NOT commit terraform.tfvars to git!

resource_group_name = "rg-restaurant-recommendation"
location          = "West Europe"
container_name    = "restaurant-api"
environment       = "dev"

# Docker image (set by CI/CD or manually)
docker_image = "ghcr.io/YOUR_USERNAME/restaurant-recommendation:latest"

# MongoDB Atlas connection (sensitive - use environment variable in CI)
# mongodb_uri = "mongodb+srv://user:pass@cluster:mongodb.net/restaurant_db"
```

2. GitHub Actions Workflows (TO DO)

2.1 Directory Structure

```
.github/
└── workflows/
    ├── ci.yml      # Test on PR
    ├── cd.yml      # Deploy on push to main
    └── validate.yml # Validate restaurants.json
```

2.2 .github/workflows/ci.yml (Test on PR)

yaml

name: CI - Test

on:

pull_request:

branches: [main]

push:

branches: [main]

jobs:

test:

name: Run Tests

runs-on: ubuntu-latest

steps:

- name: Checkout code

uses: actions/checkout@v4

- name: Set up Python

uses: actions/setup-python@v5

with:

python-version: '3.11'

cache: 'pip'

- name: Install dependencies

run: |

python -m pip install --upgrade pip

pip install -r requirements.txt

- name: Run unit tests

run: pytest tests/ -v --tb=short

- name: Lint with flake8 (optional)

run: |

pip install flake8

flake8 app/ --max-line-length=120 --ignore=E501

continue-on-error: true

validate-data:

name: Validate restaurants.json

runs-on: ubuntu-latest

steps:

- name: Checkout code

uses: actions/checkout@v4

- name: Set up Python

uses: actions/setup-python@v5

with:

with:

python-version: '3.11'

- **name:** Install dependencies

run: pip install pydantic

- **name:** Validate JSON schema

run: |

 python -c "

 import json

 from pathlib import Path

Load and validate JSON syntax

```
data = json.loads(Path('restaurants/restaurants.json').read_text())
```

Check it's a list

```
assert isinstance(data, list), 'Expected array'
```

Check each entry has required fields

```
required = {'name', 'style', 'address', 'vegetarian', 'openHour', 'closeHour'}
```

```
for i, entry in enumerate(data):
```

```
  fields = set(entry.keys())
```

```
  missing = required - fields
```

```
  extra = fields - required
```

```
  if missing or extra:
```

```
    print(f'Entry {i}: missing={missing}, extra={extra}')
```

Note: We skip invalid, don't fail (per spec)

```
print(f'Validated {len(data)} entries')
```

"

docker-build:

name: Build Docker Image

runs-on: ubuntu-latest

needs: [test]

steps:

- **name:** Checkout code

uses: actions/checkout@v4

- **name:** Set up Docker Buildx

uses: docker/setup-buildx-action@v3

- **name:** Build image (test only)

uses: docker/build-push-action@v5

with:

context: .

push: false

tags: restaurant-recommendation:test

cache-from: type=gha

cache-to-type=grid

cache-to: type=gha,mode=max

2.3 .github/workflows/cd.yml (Deploy on Push to Main)

```
yaml
```

```
name: CD - Deploy

on:
  push:
    branches: [main]
  workflow_dispatch: # Allow manual trigger

env:
  REGISTRY: ghcr.io
  IMAGE_NAME: ${{ github.repository }}

jobs:
  build-and-push:
    name: Build & Push Docker Image
    runs-on: ubuntu-latest
    permissions:
      contents: read
      packages: write

    outputs:
      image_tag: ${{ steps.meta.outputs.tags }}

    steps:
      - name: Checkout code
        uses: actions/checkout@v4

      - name: Set up Docker Buildx
        uses: docker/setup-buildx-action@v3

      - name: Log in to GitHub Container Registry
        uses: docker/login-action@v3
        with:
          registry: ${{ env.REGISTRY }}
          username: ${{ github.actor }}
          password: ${{ secrets.GITHUB_TOKEN }}

      - name: Extract metadata
        id: meta
        uses: docker/metadata-action@v5
        with:
          images: ${{ env.REGISTRY }}/{{ env.IMAGE_NAME }}
          tags: |
            type=sha,prefix=
            type=raw,value=latest

      - name: Build and push
        uses: docker/build-push-action@v5
        with:
          tag: ${{ steps.meta.outputs.image_tag }}
```

```
with:  
  context: .  
  push: true  
  tags: ${{ steps.meta.outputs.tags }}  
  labels: ${{ steps.meta.outputs.labels }}  
  cache-from: type=gha  
  cache-to: type=gha,mode=max  
  
deploy:  
  name: Deploy to Azure  
  runs-on: ubuntu-latest  
  needs: [build-and-push]  
  environment: production  
  
steps:  
  - name: Checkout code  
    uses: actions/checkout@v4  
  
  - name: Setup Terraform  
    uses: hashicorp/setup-terraform@v3  
    with:  
      terraform_version: 1.6.0  
  
  - name: Azure Login  
    uses: azure/login@v1  
    with:  
      creds: ${{ secrets.AZURE_CREDENTIALS }}  
  
  - name: Terraform Init  
    working-directory: terraform  
    run: terraform init  
  
  - name: Terraform Plan  
    working-directory: terraform  
    run: |  
      terraform plan \  
      -var="docker_image=${{ env.REGISTRY }}/${{ env.IMAGE_NAME }}:latest" \  
      -var="mongodb_uri=${{ secrets.MONGODB_URI }}" \  
      -out=tfplan  
  
  - name: Terraform Apply  
    working-directory: terraform  
    run: terraform apply -auto-approve tfplan  
  
  - name: Output API URL  
    working-directory: terraform  
    run: |  
      echo "API URL: $(terraform output -raw api_url)"  
      echo "REST Endpoint: $(terraform output -raw rest_endpoint)"
```

```
smoke-test:  
  name: Smoke Test  
  runs-on: ubuntu-latest  
  needs: [deploy]  
  
steps:  
  - name: Wait for deployment  
    run: sleep 60  
  
  - name: Test health endpoint  
    run: |  
      API_URL="${{ needs.deploy.outputs.api_url }}"  
      curl -f "${API_URL}/health" || exit 1  
  
  - name: Test REST endpoint  
    run: |  
      API_URL="${{ needs.deploy.outputs.api_url }}"  
      curl -f "${API_URL}/rest?query=italian" || exit 1
```

2.4 .github/workflows/validate.yml (Optional - Dedicated Validation)

yaml

name: Validate Data

on:

push:

paths:

- 'restaurants/restaurants.json'

pull_request:

paths:

- 'restaurants/restaurants.json'

jobs:

validate:

name: Validate restaurants.json

runs-on: ubuntu-latest

steps:

- name: Checkout code

uses: actions/checkout@v4

- name: Set up Python

uses: actions/setup-python@v5

with:

python-version: '3.11'

- name: Install dependencies

run: pip install pydantic

- name: Run validation

run: |

 python -c "

 import json

 import sys

```
required = {'name', 'style', 'address', 'vegetarian', 'openHour', 'closeHour'}
```

```
valid_veg = {'yes', 'no'}
```

with open('restaurants/restaurants.json') as f:

```
  data = json.load(f)
```

```
invalid_count = 0
```

for i, entry in enumerate(data):

```
  errors = []
```

Check fields

```
  fields = set(entry.keys())
```

 if missing := required - fields:

```
    errors.append(f'missing: {missing}')
```

 if extra := fields - required:

```

if extra := fields - required:
    errors.append(f'extra: {extra}')

# Check vegetarian value
if entry.get('vegetarian', '').lower() not in valid_veg:
    errors.append(f'vegetarian must be yes/no')

if errors:
    print(f'Entry {i} ({entry.get("name", "unknown")}): {errors}')
    invalid_count += 1

print(f'\nTotal: {len(data)} entries, {invalid_count} invalid')

# Note: We don't fail on invalid entries (per spec: skip invalid rows)
if invalid_count > 0:
    print('WARNING: Some entries are invalid and will be skipped during loading')
"
```

3. Required GitHub Secrets

Configure these in GitHub Repository Settings → Secrets and Variables → Actions:

Secret Name	Description	How to Get
MONGODB_URI	MongoDB Atlas connection string	Atlas → Connect → Drivers → Copy
AZURE_CREDENTIALS	Azure service principal JSON	See below
AZURE_SUBSCRIPTION_ID	Azure subscription ID	Azure Portal

Creating Azure Service Principal

bash

```

# Create service principal with Contributor role
az ad sp create-for-rbac \
--name "github-restaurant-app" \
--role Contributor \
--scopes /subscriptions/<SUBSCRIPTION_ID> \
--sdk-auth

# This outputs JSON - copy entire output to AZURE_CREDENTIALS secret

```

Alternative: Individual Azure Secrets

Instead of `AZURE_CREDENTIALS`, you can use:

- `ARM_CLIENT_ID`
- `ARM_CLIENT_SECRET`
- `ARM_SUBSCRIPTION_ID`
- `ARM_TENANT_ID`

4. Implementation Checklist

Phase 1: Terraform Setup

- Create `terraform/` directory
- Create `providers.tf` with Azure provider
- Create `variables.tf` with all inputs
- Create `main.tf` with resource group and container
- Create `outputs.tf` with API URL outputs
- Create `terraform.tfvars.example`
- Add `terraform/` to `.gitignore` for tfvars and tfstate
- Test locally with `terraform plan`

Phase 2: GitHub Actions

- Create `.github/workflows/` directory
- Create `ci.yml` for PR testing
- Create `cd.yml` for deployment
- Configure GitHub Secrets
- Test workflow with a PR

Phase 3: Integration Testing

- Push to main to trigger deployment
- Verify container is running in Azure
- Test `/health` endpoint
- Test `/rest?query=...` endpoint
- Verify logging works

Phase 4: Documentation

- Update README with deployment instructions
- Document GitHub Secrets requirements
- Add architecture diagram with Azure

5. File Updates Needed

`.gitignore` additions

```
gitignore
```

```
# Terraform
terraform/.terraform/
terraform/*.tfstate
terraform/*.tfstate.*
terraform/.terraform.lock.hcl
terraform/terraform.tfvars
terraform/tfplan
```

```
# Local environment
.env
*.pem
*.key
```

README.md additions

Add section for production deployment:

```
markdown
```

Production Deployment

Prerequisites

- Azure subscription
- GitHub repository with Actions enabled
- MongoDB Atlas cluster

GitHub Secrets Required

- `MONGODB_URI`: MongoDB Atlas connection string
- `AZURE_CREDENTIALS`: Azure service principal JSON

Deploy

1. Push to `main` branch
2. GitHub Actions builds and pushes Docker image
3. Terraform applies infrastructure changes
4. Container Instance starts with new image

Manual Terraform

```
\```
cd terraform
terraform init
terraform plan -var="mongodb_uri=$MONGODB_URI" -var="docker_image=ghcr.io/user/repo:latest"
terraform apply
\````
```

6. Estimated Effort

Task	Estimated Time
Terraform files	2-3 hours
GitHub Actions workflows	2-3 hours
Testing & debugging	2-4 hours
Documentation updates	1 hour
Total	7-11 hours

7. Questions to Confirm Before Starting

1. **Azure Region:** Which region to deploy? (affects latency)
2. **Container Size:** 1 CPU / 1.5GB enough?
3. **Domain:** Custom domain or Azure-provided FQDN?
4. **SSL/TLS:** Required for production?
5. **Scaling:** Single instance or multiple?
6. **Monitoring:** Azure Monitor integration needed?